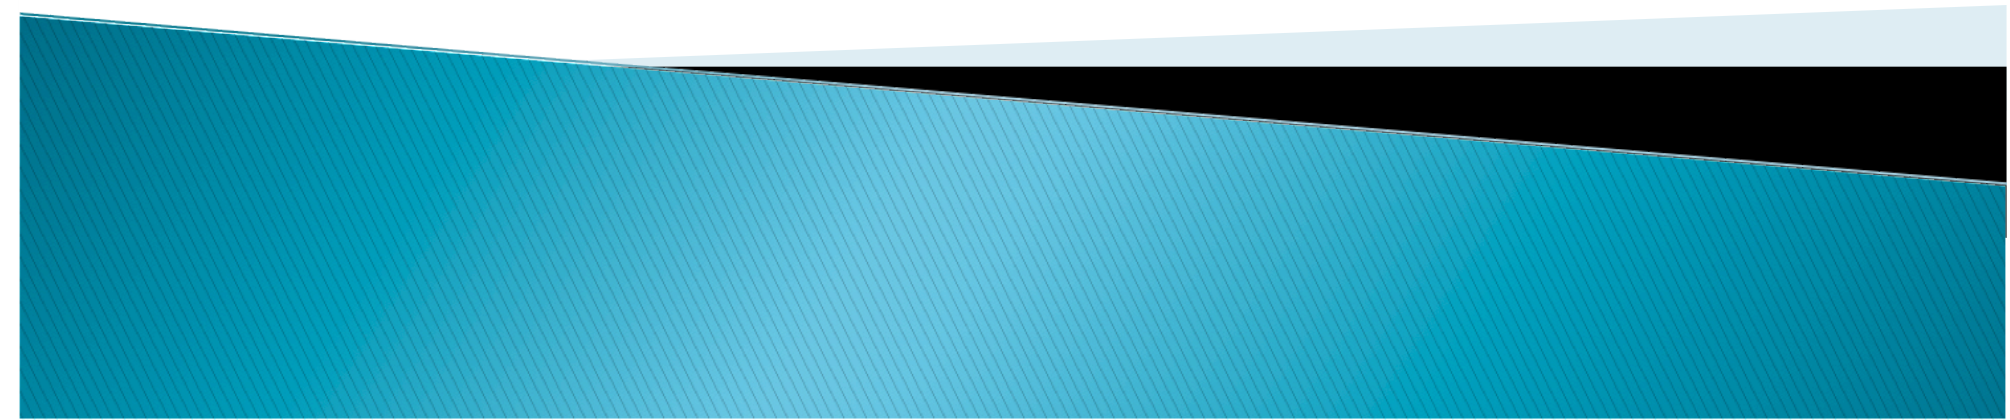
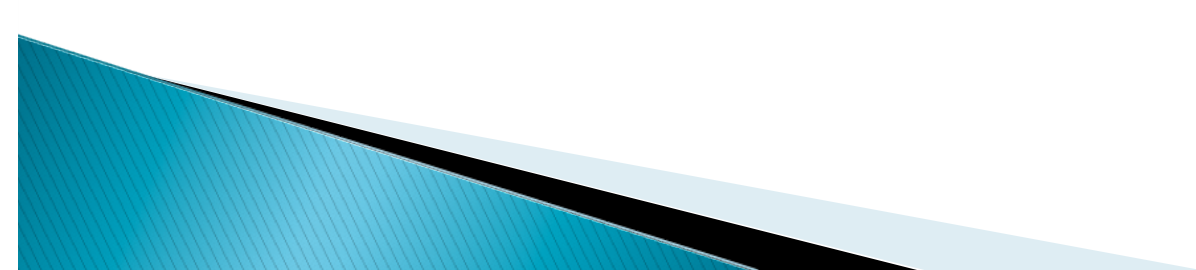


Chapter 5: CPU Scheduling



Shortest-Job-First (SJF) Scheduling


- Associate with each process the length of its next CPU burst. Use these lengths to schedule the process with the shortest time
- Two schemes:
 - **nonpreemptive** – once CPU given to the process it cannot be preempted until completes its CPU burst
 - **preemptive** – if a new process arrives with CPU burst length less than remaining time of current executing process, preempt. This scheme is known as the Shortest-Remaining-Time-First (SRTF)
- SJF is optimal – gives minimum average waiting time for a given set of processes



Example of Non-Preemptive SJF

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	7
P_2	2	4
P_3	4	1
P_4	5	4

SJF (non-preemptive)




- Waiting times $p1=0$, $p2=6$, $p3=3$, $p4=7$
- Average waiting time = $(0 + 6 + 3 + 7)/4 = 4$

Example of Preemptive SJF

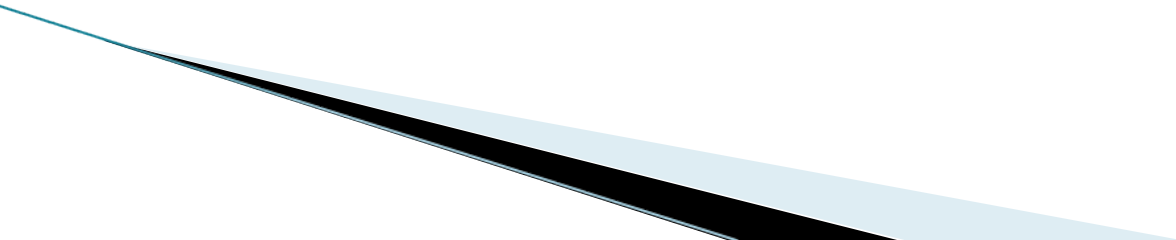
<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	7
P_2	2	4
P_3	4	1
P_4	5	4

SJF (preemptive)



- Waiting times $p1=9$, $p2=1$, $p3=0$, $p4=2$
- Average waiting time = $(9 + 1 + 0 + 2)/4 = 3$


Priority Scheduling

- A priority number (integer) is associated with each process
 - The CPU is allocated to the process with the highest priority (smallest integer □ highest priority)
 - Preemptive
 - nonpreemptive
 - SJF is a priority scheduling where priority is the predicted next CPU burst time
 - **Problem** □ Starvation – low priority processes may never execute
 - **Solution** □ Aging – as time progresses increase the priority of the process
- 

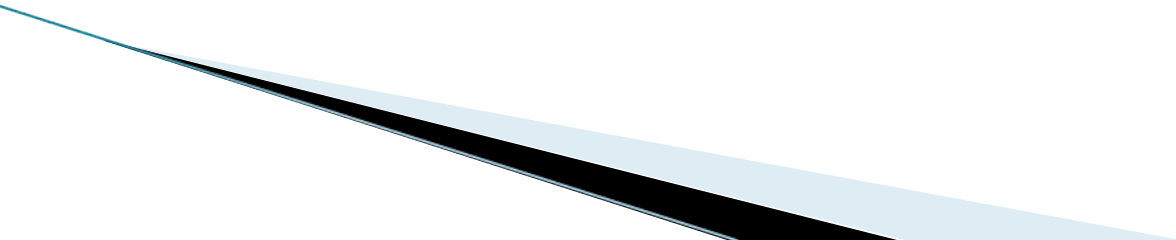
Example of Non-Preemptive Priority scheduling

<u>Process</u>	<u>Burst Time</u>	<u>Priority</u>
P_1	10	3
P_2	1	1
P_3	2	4
P_4	1	5
P_5	5	2

- Waiting times $p_1=6, p_2=0, p_3=16, p_4=18, p_5=1$
- Average waiting time = $(6 + 0 + 16 + 18 + 1) / 5 = 8.2$




Round Robin (RR)

- Each process gets a small unit of CPU time (***time quantum***), usually 10-100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue.
 - If there are n processes in the ready queue and the time quantum is q , then each process gets $1/n$ of the CPU time in chunks of at most q time units at once. No process waits more than $(n-1)q$ time units.
- 

Example of RR with Time Quantum = 20

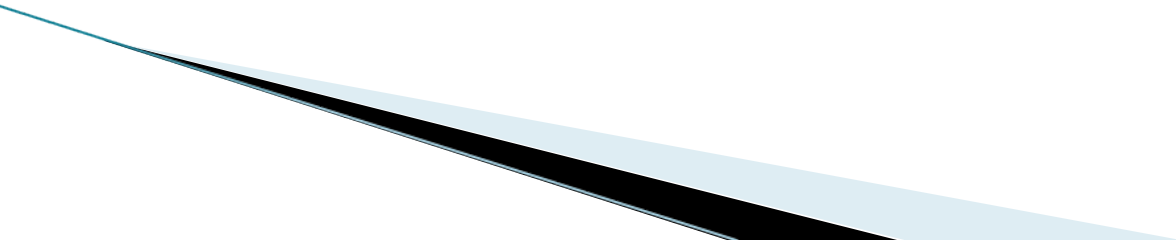
<u>Process</u>	<u>Burst Time</u>
P_1	53
P_2	17
P_3	68
P_4	24

- The Gantt chart is:



- Waiting times $p_1=81, p_2=20, p_3=94, p_4=97$
- Average waiting time = 73

Round Robin (RR)

- Performance of the RR algorithm depends heavily on the size of the time quantum.
 - If the time **quantum is extremely large**, the RR policy is the same as the FCFS policy
 - If the **time quantum is extremely small** (say, 1 millisecond), the RR approach is called processor sharing and creates the appearance that each of n processes has its own processor running at $1/n$ the speed of the real processor. (context switching overhead)
- 

Time Quantum and Context Switch Time

