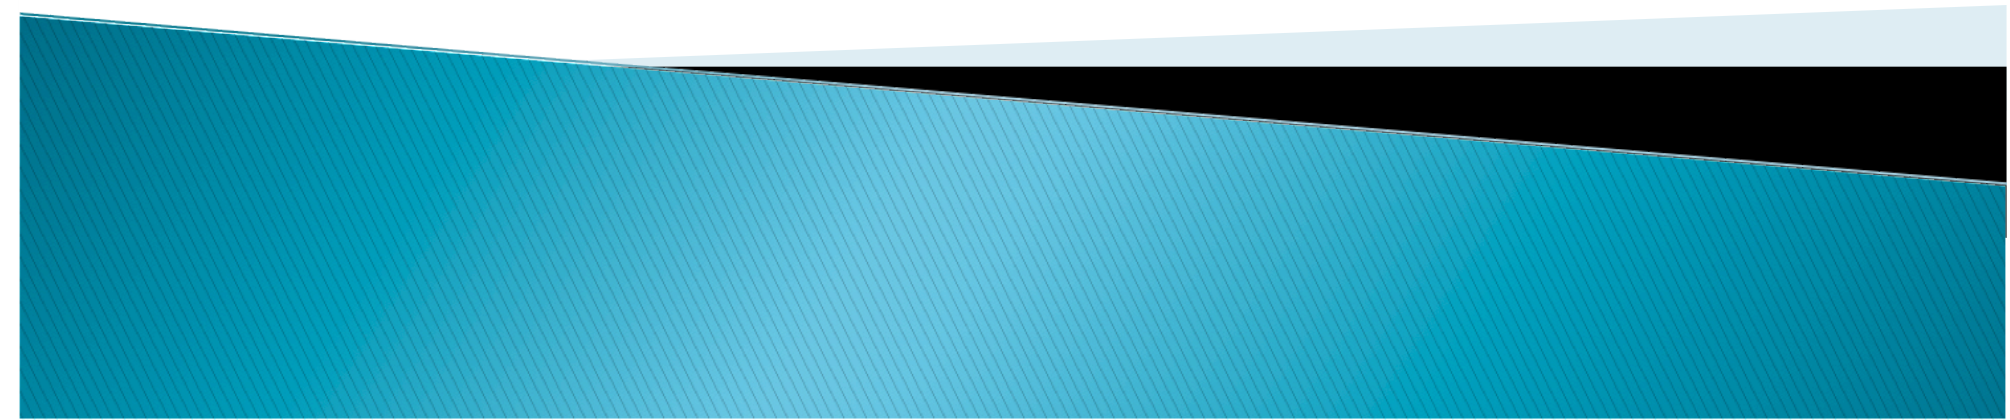


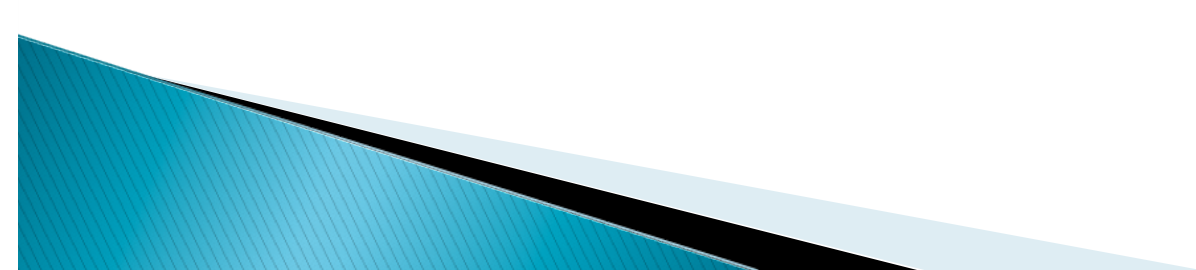
# Chapter 6: Process Synchronization

---



# Synchronization Hardware

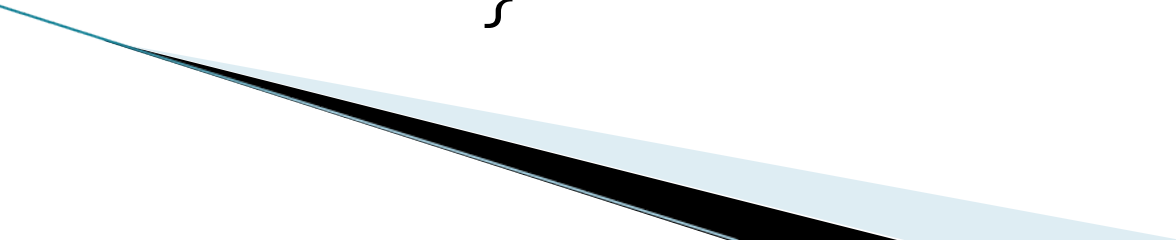
- Previously we studied a software based solution to critical section problem (Peterson solution).
- Many systems provide hardware support for critical section code.
- **Uniprocessors** – We could disable interrupts for preemption so when a shared data is being modified no preemption will occur.
  - Generally too inefficient on **multiprocessor** systems. Disabling interrupts on a multiprocessor is very time consuming, as the message is passed to all the processors
- Modern machines provide special atomic hardware instructions
  - Atomic = non-interruptable



# TestAndSet() instruction

- The important characteristic is that this instruction is executed atomically.
- If two TestAndSet C) instructions are executed simultaneously (each on a different CPU), they will be executed sequentially in some arbitrary order.

```
boolean TestAndSet (boolean *target)
{
    boolean rv = *target;
    *target = TRUE;
    return rv;
}
```



# TestAndSet() instruction (cntd)

- If the machine supports the TestAndSet () instruction, then we can implement mutual exclusion by declaring a global Boolean variable lock, initialized to false.
- Solution using TestAndSet

- Shared boolean variable lock is initialized to false.

```
do {  
    while ( TestAndSet (&lock ))  
        ; /* do nothing  
  
        // critical section  
  
    lock = FALSE;  
  
        // remainder section  
  
} while ( TRUE);
```

# Semaphore

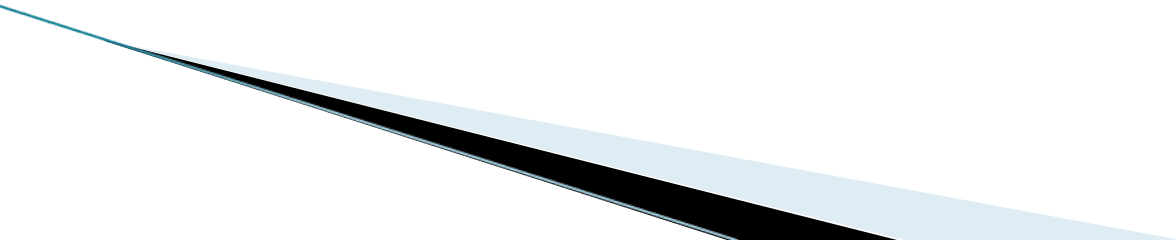
Another synchronization tool is called semaphore. A semaphore  $S$  is an integer variable that is accessed only through two standard *atomic operations*: **wait ()** and **signal ()**.

- Wait is represented by **P()** and signal by **V()**.

```
wait (S) {  
    while S <= 0  
        ; // no-op  
    S--;  
}
```

```
signal (S) {  
    S++;  
}
```

# Semaphore (cntd)

- **Binary semaphore** – integer value can range only between 0 and 1; can be simpler to implement
    - Also known as mutex locks
  
  - **Counting semaphore** – integer value can range over an unrestricted domain.
    - Used to control access to a given resource consisting of a finite number of instances.
    - Semaphore is initialized to the number of resources available. Each process performs a waitQ operation on the semaphore (thereby decrementing the count).
    - When a process releases a resource, it performs a signal () operation (incrementing the count).
    - When the count for the semaphore goes to 0, all resources are being used. After that, processes that wish to use a resource will block until the count becomes greater than 0.
- 

# Semaphore (cntd)

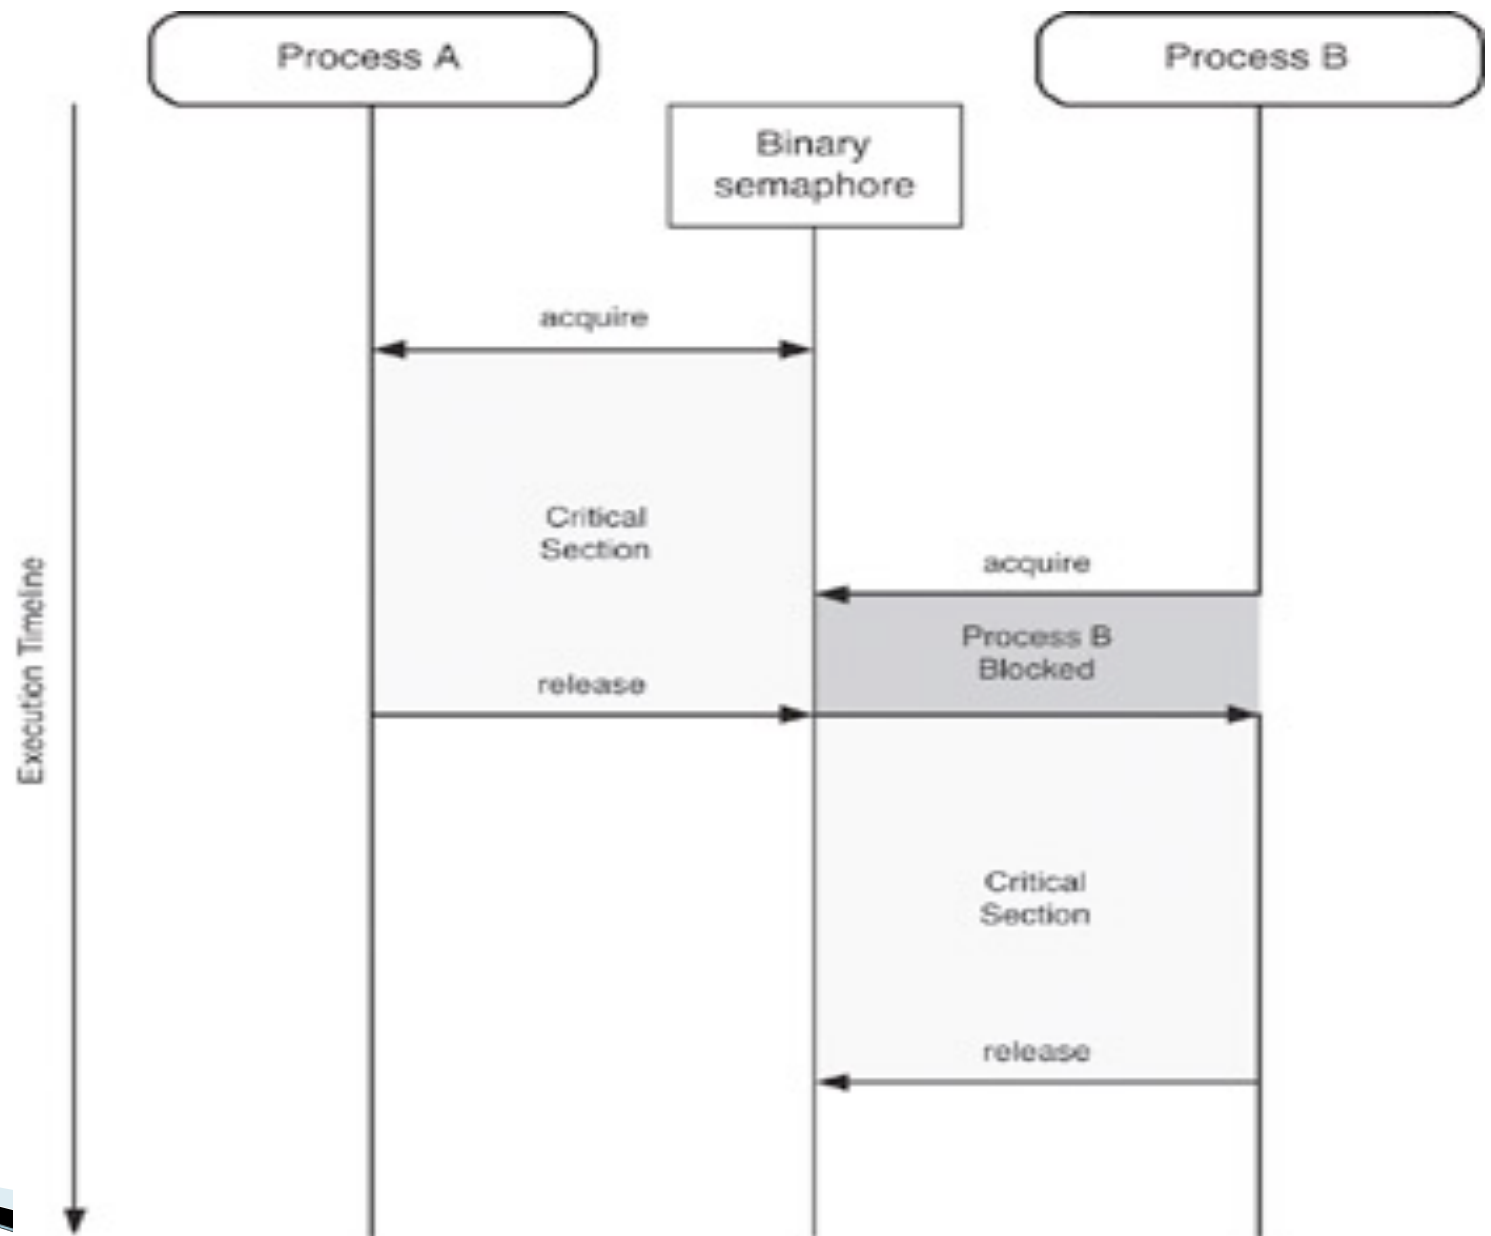
- **Mutual Exclusion**

- We can use binary semaphores to deal with the critical-section problem for multiple processes. The  $n$  processes share a binary semaphore (mutex lock) initialized to 1.

- Each process  $P_i$  is organized as shown

- Semaphore  $S$ ; // initialized to 1
- wait ( $S$ );  
    Critical Section  
    signal ( $S$ );

# Semaphore (cntd)

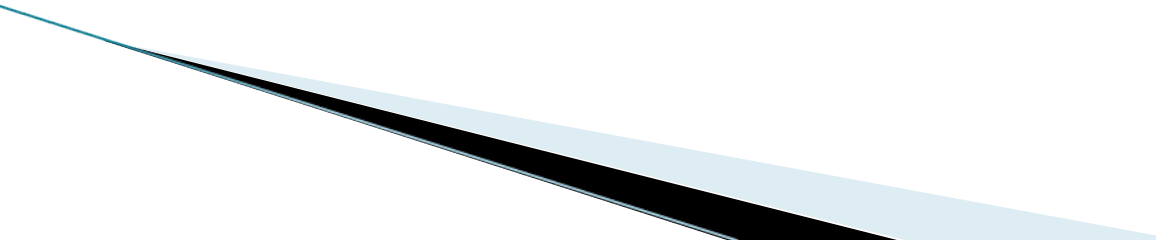




# Semaphore (cntd)

- Suppose we have 3 printers shared between 5 processes. Write code using semaphore to provide synchronized access to the printers.

# Semaphore (cntd)

- **Implementation**– Main disadvantage of the semaphore definition given previously is that every process waiting for semaphore loops continuously in the entry code.
  - Looping wastes CPU cycles that some other process might be able to use productively. This type of semaphore is also called a **spinlock** (process "spins" while waiting for the lock).
  - To eliminate busy waiting when a process has to wait for semaphore it is *blocked*.
  - The block operation places a process into a waiting queue associated with the semaphore. Hence a waiting process is not scheduled so no cpu cycles are wasted.
- 

# Semaphore (cntd)

- **Implementation**– semaphores under this definition is defined as a "C" struct:

```
typedef struct {  
    int value;  
    struct process *list;  
} semaphore;
```

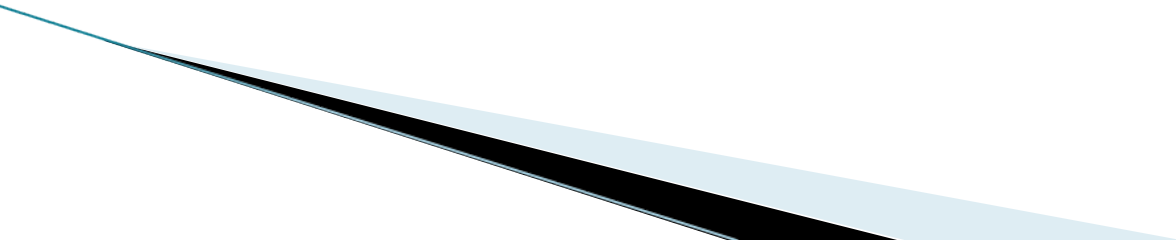
- **Implementation of wait:**

```
wait (S){  
    S->value--;  
    if (S->value < 0) {  
        add this process to waiting queue  
        block();  
    }  
}
```

- Implementation of signal:**

```
Signal (S){  
    S->value++;  
    if (S->value <= 0) {  
        remove a process P from the  
        waiting queue  
        wakeup(P);  
    }  
}
```

# Semaphore (cntd)

- The `block()` operation suspends the process that invokes it.
  - The `wakeup(P)` operation resumes the execution of a blocked process `P`. These two operations are provided by the operating system as basic system calls.
  - Note that, although under the classical definition of semaphores with busy waiting the semaphore value is never negative, this implementation may have negative semaphore values.
  - If the semaphore value is negative, its magnitude is the number of processes waiting on that semaphore.
- 

# Deadlock and Starvation

- **Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes
- Let **S** and **Q** be two semaphores initialized to 1

*P0*

wait (S);  
wait (Q);

.

.

.

signal (S);  
signal (Q);

*P1*

wait (Q);  
wait (S);

.

.

.

signal (Q);  
signal (S);

- **Starvation** – indefinite blocking. A process may never be removed from the semaphore queue in which it is suspended.